

Competency category -->

This document explains how the AI Interview Coach Bot uses 21 competency categories to organize and evaluate interview answers efficiently. The system processes 1,470 reference answers from a CSV file, categorizing them by skills like Leadership, Communication, Technical Expertise, and others. When a user answers a question, the system extracts the relevant competency (e.g., "Leadership") and searches only within that category's 100 answers instead of all 1,470, making the search 15x faster (3ms vs ~50ms) and more accurate. The competency-based filtering improves matching relevance, ensures users are compared against similar skill demonstrations, and provides transparency by showing which skills are being evaluated. This architecture is integrated throughout the codebase—from organizing data in `reference_answer_loader.py` to filtering searches for technical and behavioral questions in `main.py`—resulting in better performance, higher accuracy (85% correlation with human scores), and a superior user experience.

Csv to dictionay-->

This document explains the complete data flow of how the AI Interview Coach Bot transforms CSV data into a searchable dictionary structure and uses it to predict answer quality. The system starts with a CSV file containing 1,470 interview question-answer pairs with human scores, loads it into a pandas DataFrame, then converts it into a nested dictionary organized by 21 competency categories (e.g., "Leadership" → 100 answers, "Communication" → 70 answers). When a user answers a question, the system finds the best matching reference answer by searching within the relevant competency category using keyword overlap, then performs a multi-way comparison using TF-IDF cosine similarity (50% weight), keyword overlap (30% weight), and length ratio (20% weight) to calculate a combined reference score (0-5 points), which is added to length and relevance scores for a final prediction out of 10. This dictionary-based approach provides O(1) lookup speed, 85% correlation with human scores, and eliminates the need for repeated CSV file I/O, making the evaluation both fast and accurate.

Dataset-->

This comprehensive documentation details the 4 primary datasets powering the AI Interview Coach Bot, which collectively contain 1,582 interview questions/Q&A pairs across technical and behavioral categories. The datasets are divided into two evaluation formats: Questions-Only (Deep Learning with 111 questions and Web Development with 80 questions) that evaluate answers against question keywords without reference answers, and Q&A Pairs (Web Dev with 44 expert-scored pairs rated 8-10/10 and Behavioral with 1,470 STAR-format examples organized into 21 competency categories). The Questions-Only datasets come from Kaggle and custom curation, while the Q&A Pairs were created through expert answer writing and HR analytics data generation, with all reference answers filtered by quality scores ≥7. The system auto-detects dataset format at runtime and applies either question-based TF-IDF evaluation or multi-way comparison (50% TF-IDF similarity + 30% keyword overlap + 20% length ratio), providing both flexible concept testing and learning resources that show users what high-quality answers look like across different technologies, difficulty levels, and job competencies.

checking logic-->

This document explains how the AI Interview Coach Bot evaluates user answers using TF-IDF (Term Frequency-Inverse Document Frequency) semantic analysis rather than comparing against stored correct answers. The system loads questions from CSV files (which contain only questions, not answers), takes user input, preprocesses it using NLTK (tokenization, stop word removal, lemmatization), and calculates TF-IDF vectors to measure semantic similarity through cosine similarity. Scoring is based on three factors: length score (0-2 points for optimal word count), question relevance (0-4 points from TF-IDF similarity), and technical depth (0-4 points for explanation indicators like "because" or "therefore"), totaling up to 10 points. All answers are saved to `logs/session_log.txt` with timestamps, questions, scores, and feedback. This approach is superior to traditional keyword matching because it understands meaning rather than just detecting keywords, making it harder to game the system—for example, keyword stuffing scores poorly (1/10) while thoughtful explanations score well (9/10), whereas simple keyword matching would score them backwards.

implementation-->

This implementation summary documents the successful creation of a multi-way answer comparison system that evaluates user interview answers against 1,470 reference Q&A pairs from real interview data using three weighted metrics: TF-IDF similarity (50%), keyword overlap (30%), and length ratio (20%) to calculate a combined reference score (0-5 points) which is added to length and relevance scores for a total of 10 points. The system includes three new files (reference_answer_loader.py, test_reference_system.py, REFERENCE_ANSWER_SYSTEM.md) and enhancements to existing files (tfidf_evaluator.py, main.py) that automatically load and organize the dataset into 21 competency categories at startup, find the best matching reference answer for each question, and provide detailed feedback including TF-IDF similarity scores, keyword coverage percentages, length appropriateness, and the reference answer's human score (1-10). This improved the system's accuracy from 60% to 85%, transforming it from subjective keyword matching to objective comparison against human-scored reference answers, with comprehensive feedback that shows users exactly how their answers compare to high-quality responses across multiple dimensions.

TF/IDF-->

This document explains the TF-IDF (Term Frequency-Inverse Document Frequency) scoring system used by the AI Interview Coach Bot to evaluate answers on a 0-10 scale through three components: length score (0-2 points) that penalizes very short (<5 words) or overly long (>100 words) answers, question relevance (0-4 points) calculated using cosine similarity between TF-IDF vectors of the question and answer, and technical depth (0-4 points) based on explanatory indicators like "because," "therefore," and "involves." The system employs NLTK preprocessing with four steps—tokenization using word_tokenize(), lowercasing, removal of 179 English stop words, and lemmatization to convert words to base forms—before calculating TF-IDF vectors and measuring semantic similarity. This approach achieves 85% accuracy compared to 60% with old keyword matching because it understands meaning rather than just counting keywords, preventing keyword stuffing (scores 2/10 instead of 8/10) while properly rewarding paraphrased correct answers (7/10 instead of 4/10), with processing time of ~0.05 seconds per answer and automatic NLTK data downloading on first run.